



**NANYANG  
TECHNOLOGICAL  
UNIVERSITY**  
**SINGAPORE**



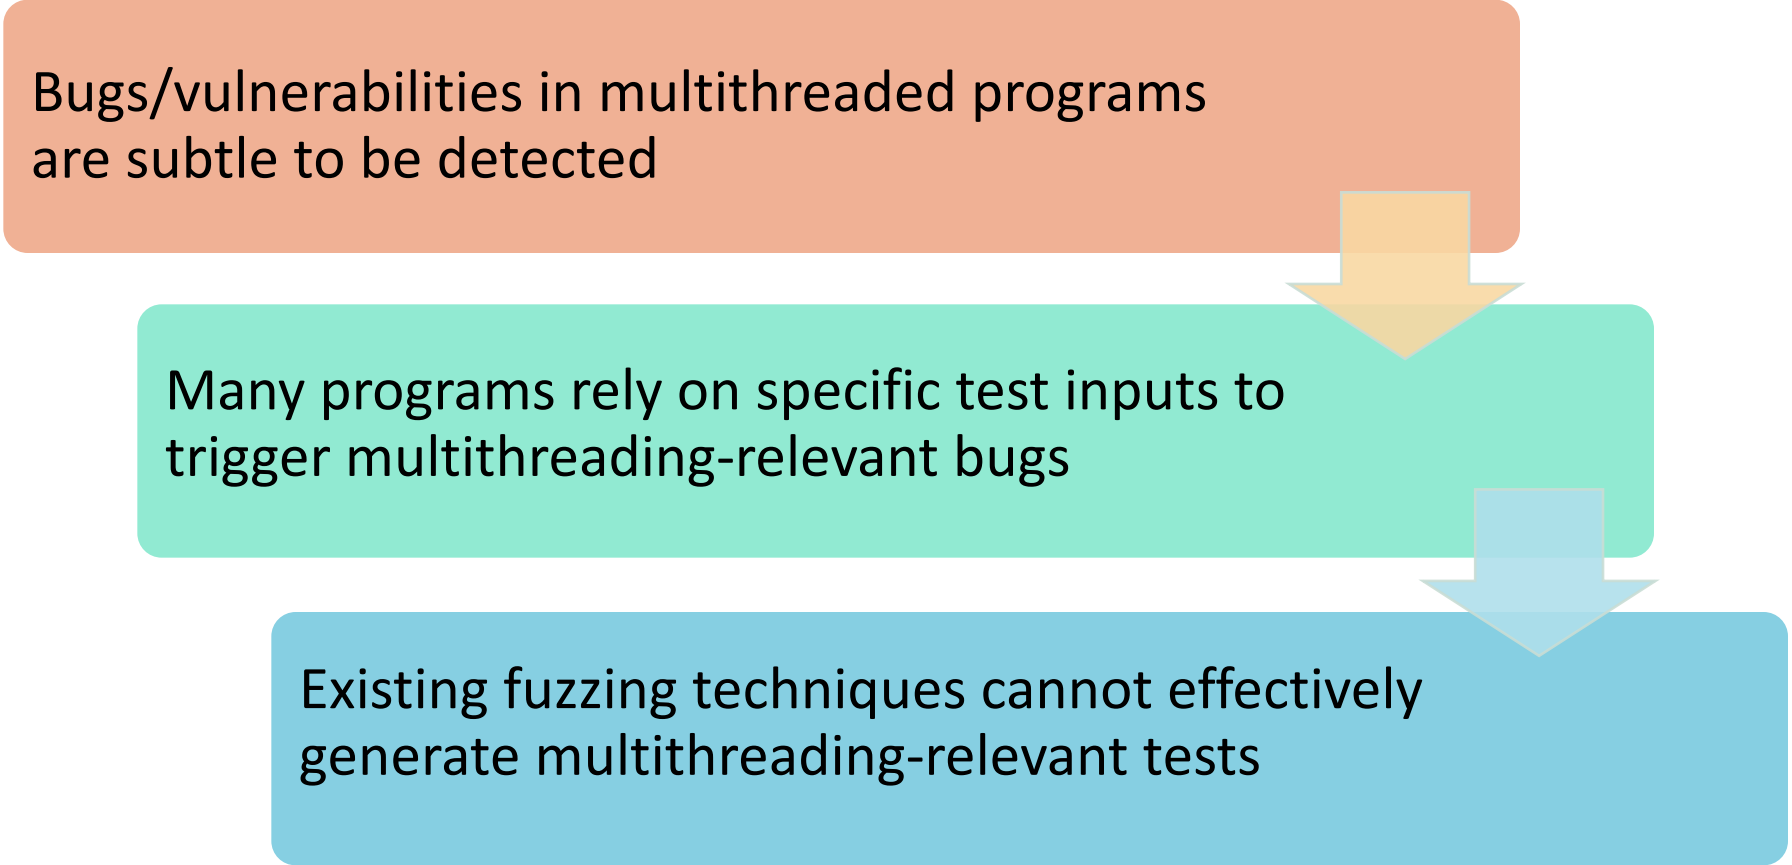
**蚂蚁金服**  
**ANT FINANCIAL**

# MUZZ: Thread-aware Grey-box Fuzzing for Effective Bug Hunting in Multithreaded Programs

**Hongxu Chen**, Shengjian Guo, Yinxing Xue, Yuelel Sui  
Cen Zhang, Yuekang Li, Haijun Wang, Yang Liu

# Background

Bugs/vulnerabilities in multithreaded programs are subtle to be detected



```
graph TD; A[Bugs/vulnerabilities in multithreaded programs are subtle to be detected] --> B[Many programs rely on specific test inputs to trigger multithreading-relevant bugs]; B --> C[Existing fuzzing techniques cannot effectively generate multithreading-relevant tests];
```

Many programs rely on specific test inputs to trigger multithreading-relevant bugs

Existing fuzzing techniques cannot effectively generate multithreading-relevant tests

# Motivation (1) – The problem

```
1  int g_var = -1;
2  void modify(int *pv) { *pv -= 2; } // ⑨
3
4  void check(char * buf) {
5      if (is_invalid(buf)) { exit(1); }
6      else { modify((int*)buf); }
7  }
8
9  char* compute(void *s_var) {
10     g_var += 1; // ①
11     g_var *= 2; // ②
12     if ((int*)s_var[0]<0) // ③
13         modify((int*)s_var); // ④
14     pthread_mutex_lock(&m); // ⑤
15     modify(&g_var); // ⑥
16     pthread_mutex_unlock(&m); // ⑦
17     return (char*)s_var; // ⑧
18 }
19
20 int main(int argc, char **argv) {
21     char * buf = read_file_content(argv[1]);
22     check(buf);
23     pthread_t T1, T2;
24     pthread_create(T1, NULL, compute, buf);
25     pthread_create(T2, NULL, compute, buf+128);
26     .....
27 }
```

- Coverage depends on test inputs

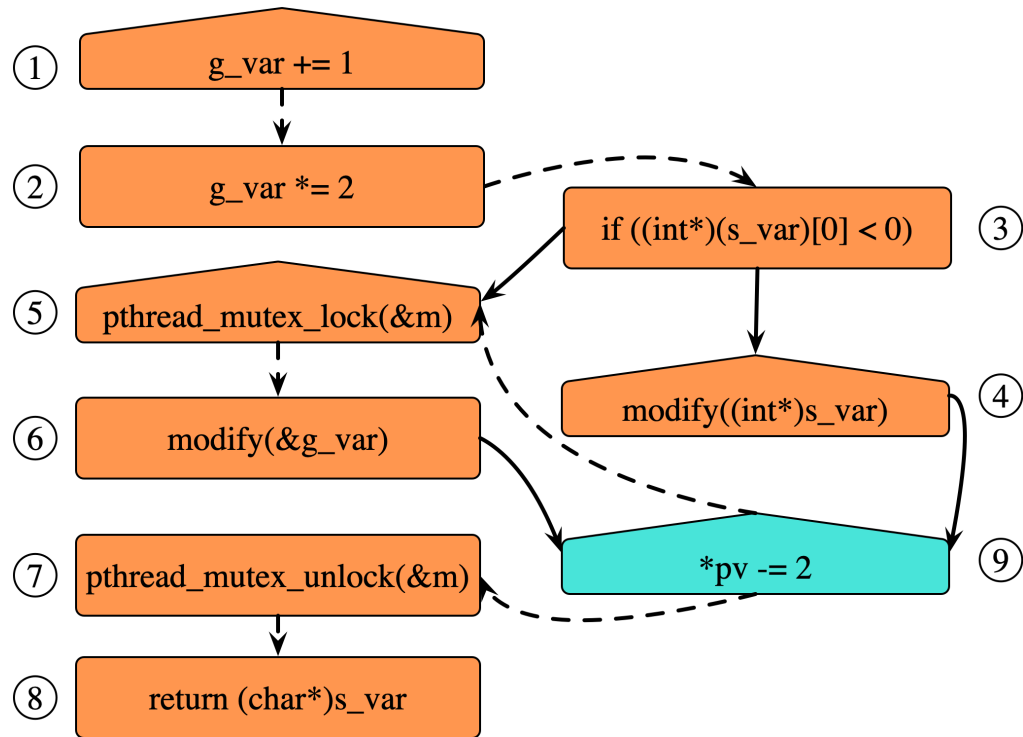
e.g., the program may or may not execute ④ according to the condition of ③, purely dependent on inputs

- Coverage depends on thread-scheduling

e.g., ① : “g\_var+=1” ② : “g\_var\*=2”

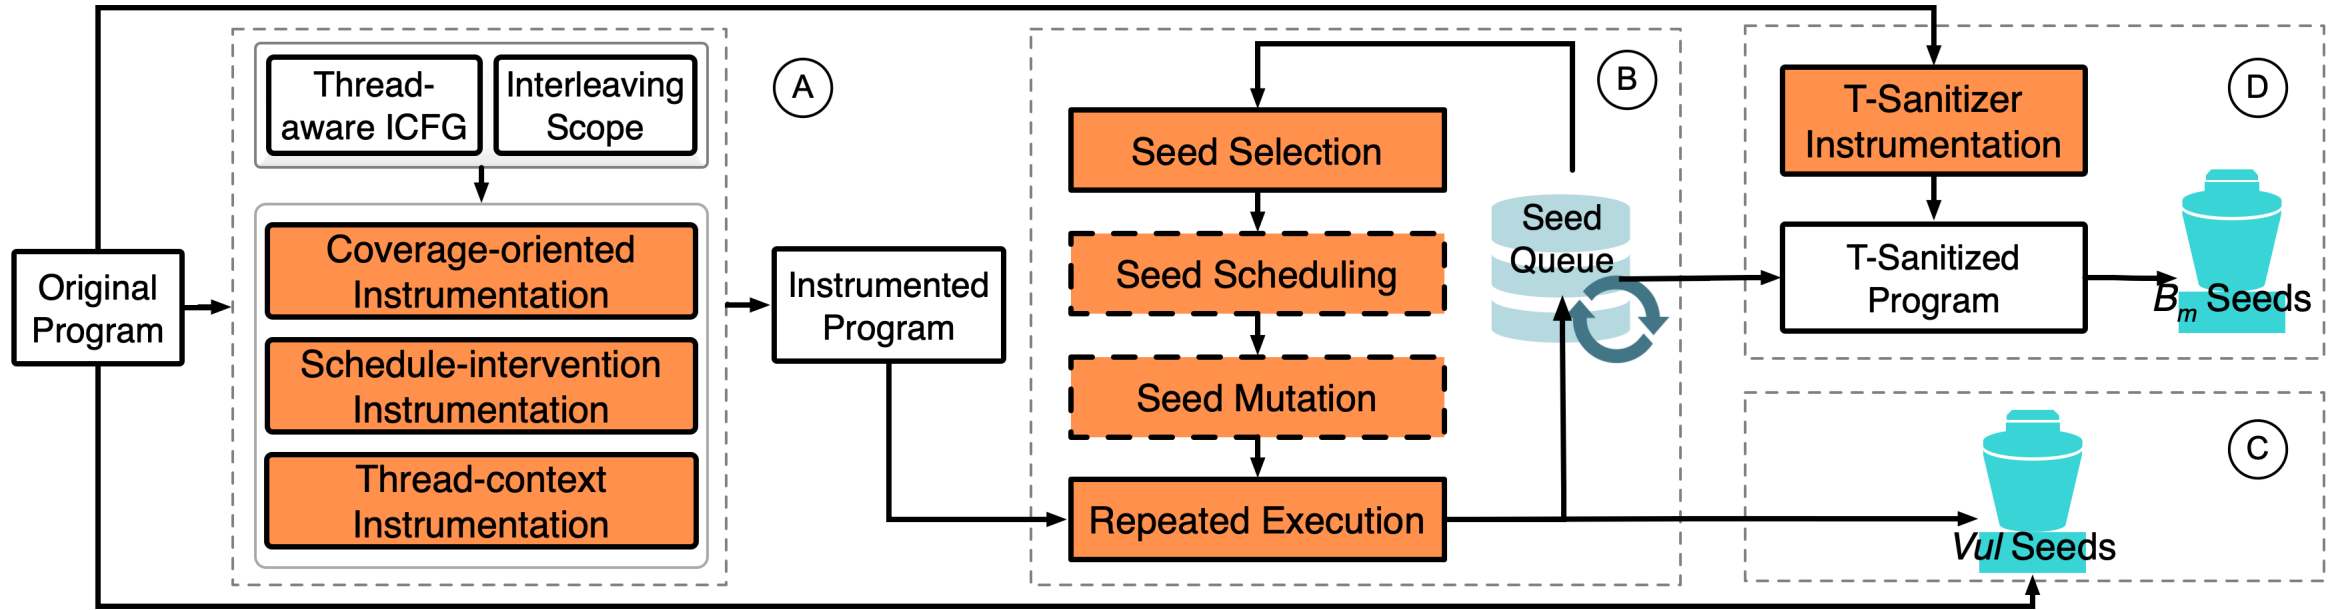
- T1: ① → T2: ① → T2: ② → T1: ② →  
g\_var=4
- T1: ① → T2: ① → T1: ② → T2: ② →  
g\_var=4
- T1: ① → T1: ② → T2: ① → T2: ② →  
g\_var=2

# Motivation (2) – Existing Solutions



- Lacking Feedback to Track Thread-interleavings and Thread-context  
e.g., ① → ①
- Lacking Schedule-intervention Across Executions
  - e.g., SAME interleaving during fuzzing?

# MUZZ Overview



Ⓐ: Static Analysis Guided Instrumentation

Ⓑ: Adaptive Dynamic Fuzzing

Ⓒ: Vulnerability Detection Analysis

Ⓓ: ThreadSanitizer Aided Concurrency-bug Revealing

# Approach (1) – Static Analysis

```
1  int g_var = -1;
2  void modify(int *pv) { *pv -= 2;} // ⑨
3
4  void check(char * buf) {
5      if (is_invalid(buf)) { exit(1); }
6      else { modify((int*)buf); }
7  }
8
9  char* compute(void *s_var) {
10     g_var += 1; // ①
11     g_var *= 2; // ②
12     if ((int*)s_var[0]<0) // ③
13         modify((int*)s_var); // ④
14     pthread_mutex_lock(&m); // ⑤
15     modify(&g_var); // ⑥
16     pthread_mutex_unlock(&m); // ⑦
17     return (char*)s_var; // ⑧
18 }
19
20 int main(int argc, char **argv) {
21     char * buf = read_file_content(argv[1]);
22     check(buf);
23     pthread_t T1, T2;
24     pthread_create(T1, NULL, compute, buf);
25     pthread_create(T2, NULL, compute, buf+128);
26     .....
27 }
```

## Identify Suspicious Interleaving Scope ( $L_m$ )

- The statements should be executed after one of *TFork*, while *TJoin* is not encountered yet
- The statements can only be executed before the invocation of *TLock* and after the invocation of *TUnLock*
- The statements should read or write at least one of the *shared variables* by different threads

# Approach (2) – Coverage-oriented Instrumentation

- Instrument more in  $L_m$ , but with certain probabilities

$$P_e(f) = \min \left\{ \frac{E(f) - N(f) + 2}{10}, 1.0 \right\}$$

$$P_s(f) = \min\{P_e(f), P_{s0}\}$$

$$P_m(f, b) = \min \left\{ P_e(f) \cdot \frac{Nm(b)}{N(b)}, P_{m0} \right\}$$

**input** : target program  $P$ , and suspicious interleaving scope  $L_m$   
**output** : program  $P$  instrumented with M-Ins deputies

```
1 for  $f \in P$  do
2   for  $b \in f$  do
3      $L_m(b) = L_m \cap b$ ;
4     if  $L_m(b) \neq \emptyset$  then
5       for  $i \in b$  do
6         if  $is\_entry\_instr(i, b)$  then
7            $P \leftarrow instrument\_cov(P, i, 1.0)$ ;
8         else if  $i \in L_m$  then
9            $P \leftarrow instrument\_cov(P, i, P_m(f, b))$ ;
10      else
11        for  $b \in f$  do
12           $i = get\_entry\_instr(b)$ ;
13           $P \leftarrow instrument\_cov(P, i, P_s(f))$ ;
```

## Approach (3) – Two Other Instrumentations

### ➤ Threading-context Instrumentation

- Track *thread IDs* and *TLock*, *TUnlock*, *TJoin*
- *Distinguish different transitions between threads*

### ➤ Schedule-intervention Instrumentation

- Using *pthread\_setschedparam* to adjust thread priority and apply *uniformly distributed random*
- Increase thread scheduling diversities



# Approach (4) – Seed Selection

**input** : seed queue  $Q_S$ , seed  $t$  at queue front

**output** : whether  $t$  will be selected in this round

```
1 if has_new_mt_ctx( $Q_S$ ) or has_new_trace( $Q_S$ ) then
2   if cov_new_mt_ctx( $t$ ) then
3     return true;
4   else if cov_new_trace( $t$ ) then
5     return select_with_prob( $P_{ynt}$ );
6   else
7     return select_with_prob( $P_{ynn}$ );
8 else
9   return select_with_prob( $P_{nnn}$ );
```

Prioritize to select those seeds that:

- Cover new regular traces
- Cover new **thread-interleavings**

# Approach (5) – Repeated Execution

```

input : program  $\mathbb{P}_o$ , initial seed queue  $Q_S$ 
output : final seed queue  $Q_S$ , vulnerable seed files  $T_C$ 
1  $\mathbb{P}_f \leftarrow \text{instrument}(\mathbb{P}_o)$ ; // instrumentation
2  $T_C \leftarrow \emptyset$ ;
3 while True do
4    $t \leftarrow \text{select\_next\_seed}(Q_S)$ ; // seed selection
5    $M \leftarrow \text{get\_mutation\_chance}(\mathbb{P}_f, t)$ ; // seed scheduling
6   for  $i \in 1 \dots M$  do
7      $t' \leftarrow \text{mutated\_input}(t)$ ; // seed mutation
8      $\text{res} \leftarrow \text{run}(\mathbb{P}_f, t', \mathbb{N}_c)$ ; // repeated execution
9     if  $\text{is\_crash}(\text{res})$  then // seed triaging
10       $T_C \leftarrow T_C \cup \{t'\}$ ; // report vulnerable seeds
11     else if  $\text{cov\_new\_trace}(t', \text{res})$  then
12       $Q_S \leftarrow Q_S \oplus t'$ ; // preserve "effective" seeds

```

$$\mathbb{N}_c(t) = N_0 + N_v \cdot B_v, \quad B_v \in \{0, 1\} \quad \longrightarrow \quad \mathbb{N}_c(t) = N_0 + \min\{N_v, N_0 \cdot C_m(t)\}$$

# Statistics of Target Programs

ID	Project	Command Line Options	Binary Size	$T_{pp}$	$N_b$	$N_i$	$N_{ii}$	$\frac{N_{ii}-N_b}{N_b}$
<b>lbzip2-c</b>	lbzip2-2.5	lbzip2 -k -t -9 -z -f -n4 FILE	377K	7.1s	4010	24085	6208	54.8%
<b>pbzip2-c</b>	pbzip2-v1.1.13	pbzip2 -f -k -p4 -S16 -z FILE	312K	0.9s	2030	8345	2151	6.0%
<b>pbzip2-d</b>	pbzip2-v1.1.13	pbzip2 -f -k -p4 -S16 -d FILE	312K	0.9s	2030	8345	2151	6.0%
<b>pigz-c</b>	pigz-2.4	pigz -p 4 -c -b 32 FILE	117K	5.0s	3614	21022	5418	49.9%
<b>pxz-c</b>	pxz-4.999.9beta	pxz -c -k -T 4 -q -f -9 FILE	42K	1.2s	3907	30205	7877	101.6%
<b>xz-c</b>	XZ-5.3.1alpha	xz -9 -k -T 4 -f FILE	182K	8.4s	4892	34716	8948	82.9%
<b>gm-cnvt</b>	GraphicsMagick-1.4	gm convert -limit threads 4 FILE out.bmp	7.6M	224.4s	63539	383582	98580	55.1%
<b>im-cnvt</b>	ImageMagick-7.0.8-7	convert -limit thread 4 FILE out.bmp	19.4M	434.2s	128359	778631	200108	55.9%
<b>cwebp</b>	libwebp-1.0.2	cwebp -mt FILE -o out.webp	1.8M	56.3s	12117	134824	33112	173.3%
<b>vp8dec</b>	libvpx-v1.3.0-5589	vp8dec -t 4 -o out.y4m FILE	3.8M	431.6s	31638	368879	93400	195.2%
<b>x264</b>	x264-0.157.2966	x264 -threads=4 -o out.264 FILE	6.4M	1701.0s	38912	410453	103926	167.1%
<b>x265</b>	x265-3.0_Au+3	x265 -input FILE -pools 4 -F 2 -o	9.7M	78.3s	22992	412555	89408	288.9%

$T_{pp}$ : Preprocessing time

$N_b$ : Number of basicblocks

$N_i$ : Number of instructions

$N_{ii}$ : Number of MUZZ-instrumented instructions

# Evaluation (1) – Seed Generation

ID	MUZZ			MAFL			AFL			MOPT		
	$N_{all}$	$N_{mt}$	$\frac{N_{mt}}{N_{all}}$	$N_{all}$	$N_{mt}$	$\frac{N_{mt}}{N_{all}}$	$N_{all}$	$N_{mt}$	$\frac{N_{mt}}{N_{all}}$	$N_{all}$	$N_{mt}$	$\frac{N_{mt}}{N_{all}}$
lbzip2-c	8056	<b>5127</b>	<b>63.6%</b>	6307	3277(+1850)	52.0%(+11.7%)	5743	2464(+2663)	42.9%(+20.7%)	6033	2524(+2603)	41.8%(+21.8%)
pbzip2-c	381	<b>126</b>	<b>33.1%</b>	340	91(+35)	26.8%(+6.3%)	272	69(+57)	25.4%(+7.7%)	279	71(+55)	25.4%(+7.6%)
pbzip2-d	1997	<b>297</b>	<b>14.9%</b>	1706	119(+178)	7.0%(+7.9%)	1650	68(+229)	4.1%(+10.8%)	1623	62(+235)	3.8%(+11.1%)
pigz-c	1406	<b>1295</b>	<b>92.1%</b>	1355	1189(+106)	87.7%(+4.4%)	1298	1098(+197)	84.6%(+7.5%)	1176	982(+313)	83.5%(+8.6%)
pxz-c	7590	<b>5249</b>	<b>69.2%</b>	5637	3401(+1848)	60.3% (+8.8%)	5357	2470(+2779)	46.1% (+23.0%)	5576	2634(+2615)	47.2% (+21.9%)
xz-c	2580	<b>1098</b>	<b>42.6%</b>	2234	767(+331)	34.3%(+8.2%)	1953	581(+517)	29.7%(+12.8%)	1845	566(+532)	30.7%(+11.9%)
gm-cnvt	15333	<b>13774</b>	<b>89.8%</b>	14031	10784(+2990)	76.9%(+13.0%)	12453	8290(+5484)	66.6%(+23.3%)	12873	8956(+4818)	69.6%(20.3%)
im-cnvt	14377	<b>12987</b>	<b>90.3%</b>	12904	10610(+2377)	82.2%(+8.1%)	9935	7634(+5353)	76.8%(+76.8%)	10203	8012(+4975)	78.5%(+11.8%)
cwebp	11383	<b>7554</b>	<b>66.4%</b>	10389	6868(+686)	66.1% (+0.3%)	9754	5874(+1680)	60.2% (+6.1%)	9803	5869(+1685)	59.9%(+6.5%)
vpzdec	28892	<b>25593</b>	<b>88.6%</b>	27735	22507(+3086)	81.2%(+7.4%)	24397	18936(+6657)	77.6%(+11.0%)	27119	20896(+4697)	77.1%(11.5%)
x264	15138	<b>14611</b>	<b>96.5%</b>	14672	13413(+1198)	91.4% (+5.1%)	13211	11801(+2810)	89.3% (+7.2%)	12427	11202(+3409)	90.1%(+6.4%)
x265	12965	10704	<b>82.6%</b>	13858	<b>10890</b> (-186)	78.6% (+4.0%)	12980	9957(+747)	76.7% (+5.9%)	13142	10154 (+550)	77.3%(+5.3%)

MUZZ has advantages in increasing the **number** and **percentages** of **multithreading-relevant seeds** for multithreaded programs

## Evaluation (2) – Vulnerability Detection

ID	MUZZ					MAFL					AFL					MOPT				
	$N_c$	$N_c^m$	$N_v^m$	$N_c^s$	$N_v^s$	$N_c$	$N_c^m$	$N_v^m$	$N_c^s$	$N_v^s$	$N_c$	$N_c^m$	$N_v^m$	$N_c^s$	$N_v^s$	$N_c$	$N_c^m$	$N_v^m$	$N_c^s$	$N_v^s$
pbzip2-c	6	<b>6</b>	<b>1</b>	0	0	6	0(+6)	<b>1(0)</b>	0	0	0	0(+6)	0(+1)	0	0	0	0(+6)	0(+1)	0	0
pbzip2-d	15	<b>15</b>	<b>1</b>	0	0	0	0(+15)	0(+1)	0	0	0	0(+15)	0(+1)	0	0	0	0(+15)	0(+1)	0	0
im-cnvt	87	<b>63</b>	<b>4</b>	24	1	49	23(+40)	2(+2)	26	1	29	6(+57)	2(+2)	23	1	32	6(+57)	2(+2)	26	1
cwebp	19	0	0	19	1	27	0(0)	0(0)	27	1	14	0(0)	0(0)	14	1	15	0(0)	0(0)	15	1
vpdec	523	<b>347</b>	<b>2</b>	176	2	495	279(+68)	1(+1)	216	2	393	205(+142)	1(+1)	188	2	501	301(+46)	1(+1)	200	2
x264	103	<b>103</b>	<b>1</b>	0	0	88	88(+15)	<b>1(0)</b>	0	0	78	78(+25)	<b>1(0)</b>	0	0	66	66(+37)	<b>1(0)</b>	0	0
x265	43	0	0	43	1	52	0(0)	0(0)	52	1	62	0(0)	0(0)	62	1	59	0(0)	0(0)	59	1

MUZZ demonstrates superiority in exercising more multithreading-relevant crashing states and detecting **concurrency-vulnerabilities**

# Evaluation (3) – Concurrency-bug Revealing

ID	$\mathbb{P}1$								$\mathbb{P}2$							
	MUZZ		MAFL		AFL		MOPT		MUZZ		MAFL		AFL		MOPT	
	$N_e^m$	$N_B^m$	$N_e^m$	$N_B^m$	$N_e^m$	$N_B^m$	$N_e^m$	$N_B^m$	$N_e^m$	$N_B^m$	$N_e^m$	$N_B^m$	$N_e^m$	$N_B^m$	$N_e^m$	$N_B^m$
lbzip2-c	<u>469</u>	<u>1</u>	447	<u>1</u>	386	<u>1</u>	435	<u>1</u>	<b>493</b>	<b>1</b>	483	<b>1</b>	421	<b>1</b>	458	<b>1</b>
pigz-c	793	<u>1</u>	<u>803</u>	<u>1</u>	764	<u>1</u>	789	<u>1</u>	856	<b>1</b>	<b>862</b>	<b>1</b>	727	<b>1</b>	742	<b>1</b>
gm-cnvt	<u>93</u>	<u>5</u>	79	4	45	2	55	3	<b>133</b>	<b>5</b>	83	4	54	3	57	3
im-cnvt	<u>92</u>	<u>3</u>	84	<u>3</u>	58	2	56	2	<b>118</b>	<b>3</b>	117	<b>3</b>	65	2	59	2
vpxdec	<u>31</u>	<u>3</u>	17	1	23	1	18	1	<b>42</b>	<b>3</b>	22	1	25	1	22	1
x264	<u>68</u>	<u>8</u>	46	6	28	4	30	5	<b>91</b>	<b>9</b>	52	6	25	4	28	4

MUZZ outperforms competitors in revealing  
**concurrency-bugs** with fuzzer-generated seeds



# Q&A

Hongxu Chen

[hongxu.chen@ntu.edu.sg](mailto:hongxu.chen@ntu.edu.sg)

